

---

# DEMAND FORECASTING PROJECT PRESENTATION

CS613: Machine Learning  
Manjiri Khodke, Sriram Suvaidhar,  
Khushi Patel, Mokshad Sankhe

---

---

# PROBLEM STATEMENT

- The main aim is to make sure each shop receives the right amount of stock at the right time.
  - We want to avoid situations where a shop runs out of products and loses customers but also prevent sending more stock than the shop can realistically sell.
  - In simple terms, the objective is to maintain a healthy balance not too much, not too little so that shops can operate smoothly, meet customer demand, and avoid unnecessary costs or wastage.
-

---

# DATASET

- The dataset is taken from the Google Cloud public Marketplace in Big Query.
  - **Sales Data:** Daily sales transactions
    - Fields: `sales_date`, `dealer_code`, `product`, `qty`, `total_price`
  - **Stock Data:** Daily inventory snapshots
    - Fields: `benchmark_date`, `ec_id`, `product_name`, `inventory_qty`
  - **Scale:** Multiple stores × multiple products × daily timestamps
  - **Preprocessing:** Merged datasets, handled missing values, created time-series structure.
-

---

# UPGRADING STOCK FORECASTING WITH ML

- Old Method: simple average of past sales
- Limitations: Misses seasonality, trends, demand spikes, and shop differences
- New Method: ML algorithms analyzing multiple factors
- Improvement: more accurate, dynamic forecasts
- Results: better stock allocation and inventory planning
- **Reference:**


H. Malik, “A beginner’s approach to time-series with working example: Demand forecasting | Time series example with Kaggle,” Medium, Jun. 11, 2024. [Online]. Available: <https://medium.com/@humzahmalik/a-beginners-approach-to-time-series-with-working-example-c6bff9c24928>

---

---

# PROBLEM FORMULATION

- It a regression problem and not classification
- **Target Variable:** Continuous daily sales quantity (Predict how many units will sell, not just categories)
- **Why Regression?**
  - Need exact numbers for inventory ordering
  - Business decisions require precise quantities
  - Natural numeric output (0, 1, 2, ... units)



# PROGRAM ARCHITECTURE

- **Data Input:** Sales and Stock data (CSV format) are fed into the system.
  - **Data Preprocessing:** The data is merged, features are engineered, and it's split for training/testing.
  - **Modeling:** Four machine learning models (LightGBM, XGBoost, Random Forest, SVR) are trained, and an Ensemble Model is created from their weighted average.
  - **Evaluation:** The models are assessed using metrics like RMSE, WMAPE, and R<sup>2</sup> Score.
  - **Demand Forecasting:** The Ensemble Model is used to make 7-day predictions.
  - **Inventory Management:** The forecasts are processed through Inventory Rules (e.g., safety stock) to generate Status Flags (CRITICAL LOW, LOW STOCK, OK, OVERSTOCK).
  - **Output:** The system produces a Forecast Output CSV (the predictions) and Inventory Recommendations (actionable alerts) based on the status flags.

---

# **LIGHT GRADIENT BOOSTING MACHINE (LIGHTGBM)**

- Light Gradient Boosting Machine (LightGBM) serves as our primary model, selected for its superior speed and predictive performance.
  - It is a gradient boosting framework that builds trees in a leaf-wise manner, which often achieves higher accuracy compared to traditional level-wise growth.
  - The model accelerates training on large datasets through a histogram-based algorithm for efficient feature splitting and supports categorical variables natively without preprocessing.
  - Furthermore, its compatibility with GPU execution ensures scalability and production readiness for high-performance deployments.
-

---

# SUPPORT VECTOR REGRESSION (SVR)

- It is a Kernel-Based Approach for Regression.
  - This method utilizes Support Vector Regression (SVR), which adapts the principles of Support Vector Machines (SVM) from classification to predict continuous target values.
  - The approach employs the kernel trick to implicitly map input data into higher-dimensional spaces, enabling the modeling of complex, non-linear relationships. Our experimentation tested multiple configurations, including the Radial Basis Function (RBF) kernel with various regularization (C) and epsilon parameters, alongside linear and polynomial kernels for comparative analysis.
  - Given SVR's sensitivity to feature scale, all inputs were normalized using a StandardScaler prior to model training.
-

---

# COMPLEMENTARY MODELS

- **XGBoost**
    - Similar to LightGBM but uses level-wise tree growth
    - Additional regularization terms
    - Early stopping based on validation RMSE
  - **Random Forest**
    - Ensemble of decision trees
    - Bagging (bootstrap aggregation) approach
    - Built-in feature importance
-

---

# ENSEMBLE STRATEGY

- It is a Weighted Average Ensemble Approach.
- **Rationale:** Combine strengths, mitigate weaknesses
- **Weight Calculation:** Inverse of RMSE (better models get higher weight)

$$w_i = \frac{1/\text{RMSE}_i}{\sum_{j=1}^4 (1/\text{RMSE}_j)}$$

- **Final Prediction:**

$$\hat{y} = w_{\text{lgb}} \cdot \hat{y}_{\text{lgb}} + w_{\text{xgb}} \cdot \hat{y}_{\text{xgb}} + w_{\text{rf}} \cdot \hat{y}_{\text{rf}} + w_{\text{svm}} \cdot \hat{y}_{\text{svm}}$$

---

---

# FEATURE ENGINEERING STRATEGY

- The strategy focuses on Creating Predictive Signals by extracting and synthesizing information across multiple domains.
  - Temporal features are generated from dates, including day, month, quarter, and weekend indicators.
  - Historical purchase patterns are captured through lagged variables, such as 1-day and 7-day lags, and rolling averages.
  - To model store-product specificity, unique identifiers are encoded, and historical averages and standard deviations are calculated per combination.
  - Additional context is incorporated from inventory data, including ending stock levels and derived turnover ratios, alongside price signals based on average unit cost calculations.
-

---

# TRAIN-TEST METHODOLOGY

- The model development process employs a rigorous time-based splitting strategy.
  - A chronological split is critical to prevent data leakage and simulate real-world forecasting conditions.
  - Data is partitioned with an 80% training and 20% testing ratio, strictly maintaining temporal order.
  - Validation is conducted using a subset of the training data for early stopping to avoid overfitting.
  - To ensure a fair and consistent evaluation across all models, the same fixed test set is used for all performance comparisons.
-

---

# EVALUATION METRICS

- **RMSE** (Root Mean Square Error):

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

- RMSE is the main accuracy score where larger errors count more heavily toward the final result.
- **WMAPE** (Weighted Mean Absolute Percentage Error):

$$\text{WMAPE} = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{\sum_{i=1}^n y_i}$$

- WMAPE translates errors into an easy-to-understand percentage based on total sales
  - **MAE** (Mean Absolute Error):  
$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$
  - MAE gives the straightforward average of all errors, treating each mistake the same regardless of size.
-

---

# RESULTS - MODEL PERFORMANCE



| Model         | RMSE ↓  | MAE ↓  | WMAPE ↓ |
|---------------|---------|--------|---------|
| LightGBM      | 0.9573  | 0.3531 | 1.09%   |
| XGBoost       | 1.0774  | 0.4227 | 1.30%   |
| Random Forest | 2.3733  | 1.2565 | 3.86%   |
| SVM (RBF)     | 12.0908 | 6.1007 | 18.76%  |
| ENSEMBLE      | 1.0383  | 0.4559 | 1.40%   |

---




# RESULTS - FEATURE IMPORTANCE

- Analysis of our LightGBM model reveals the key drivers of predictions.
  - Top 5 Features (LightGBM Gain-based):
    - sales\_lag\_1 - Yesterday's sales (strongest predictor)
    - store\_product\_avg\_sales - Historical baseline
    - sales\_rolling\_7 - Weekly trend
    - ending\_inventory - Current stock levels
    - day\_of\_week - Weekly seasonality
  - This ranking shows that immediate recent history and store-product specific patterns are the most influential factors in our forecasts.
-

# VISUALIZATION



# VISUALIZATION



---

# FUTURE EXTENSIONS

- **External Factors Integration:**
    - Holiday calendars
    - Promotional events
    - Weather data
  - **Advanced Techniques:**
    - Deep learning (LSTM/GRU for sequence modeling)
    - Bayesian optimization for hyperparameter tuning
  - **System Improvements:**
    - Real-time updates with streaming data
    - Automated retraining pipeline
    - A/B testing framework for model updates
-

---

# CONCLUSION

- This project shows that machine learning can turn inventory management into a smarter, proactive process.
  - By using four different models together, we built a system that forecasts daily sales with high accuracy.
  - Our solution provides actionable restocking suggestions, helping businesses reduce both missed sales from stockouts and waste from overstock.
-



---

**THANK YOU! QUESTIONS?**